


1 THE PROBLEM

The reviewer is the author’s own twin




One training pipeline $\Rightarrow$ correlated blind spots. LLM judges favour output resembling their own, so the errors the author is prone to make are the ones the reviewer waves through.

The supervision trace, moreover, lives in opaque platform logs – not in an artefact a third party can replay.


2 THE PROTOCOL

Audit across vendors; put every step on the record



3 WHAT WE FOUND

The vendors disagree – not that either one is better



vertical rule = random floor from a 2000-shuffle permutation test preserving each arm’s output volume;

recall under the frozen scorer rises with finding count, so the headline 41 > 38 sits within noise once corrected.

30 synthetic increments (10 clean) · 43 seeded defects · one Constitution · three arms, identical inputs

once corrected, the order reverses	same-family	cross-vendor
strict-tier recall	38/43	37/43
chance-corrected agreement κ	0.81	0.68
false blockers on 10 clean	0	26

The cross-vendor arm returned **BLOCKED** on all thirty increments – a stub

auditor that always blocks reproduces that result exactly.

On identical inputs two vendors read one rulebook differently: the same-family auditor read CA-DATA-001 by its intent and raised nothing; the foreign auditor enforced its letter.

That decorrelation is the premise of cross-vendor audit – but this exploratory trial ($n=43$, component-level) measures **disagreement, not discrimination**. Its intended blinding was voided by a cross-vendor audit of *this repository*, whose BLOCK-grade report and our itemised dispositions are committed beside the data. · github.com/dongzhaohe321418-lab/crossaudit