


THE PROBLEM

Who audits the AI scientist?

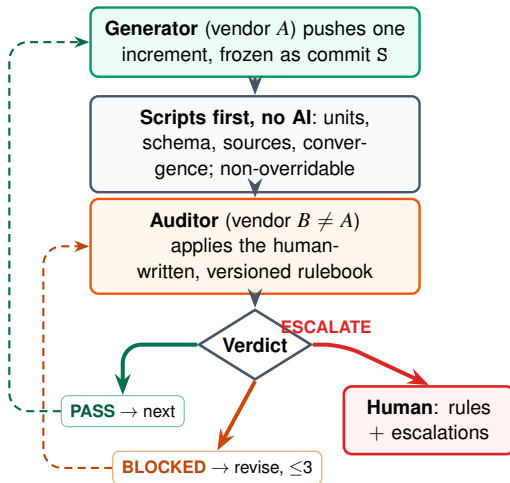


Same training pipeline $\Rightarrow$ same blind spots. LLM judges prefer output resembling their own (self-preference bias); shared errors sail through.

Internal critics in today's AI-scientist systems are the author's twin.

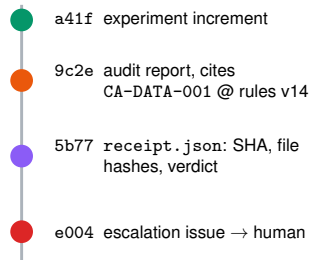
THE PROTOCOL

A second vendor, scripts first, git underneath



THE RECORD

Supervision as a public artefact



Every supervision step is a commit: replayable, diffable, citable, by anyone.

8 invariants

fork-and-run

An AI scientist should not grade its own homework. Now it doesn't have to.