

THE PROBLEM

Who audits the AI scientist?



Same training pipeline $\Rightarrow$ same blind spots:
LLM judges prefer output resembling their own
(self-preference bias); shared errors sail through review.


THE RECORD

Supervision as a public artefact

- a41f experiment increment
- 9c2e audit report: cites CA-DATA-001 @ rules v14
- 5b77 receipt.json: SHA, file hashes, verdict
- e004 escalation issue $\rightarrow$ human

Every supervision step is a commit: replayable, diffable, citable, by anyone.

THE PROTOCOL one experiment increment through the loop



Every step above is a git commit or a GitHub issue: the entire supervision history is public and replayable by anyone.

An AI scientist should not grade its own homework. Now it doesn't have to.

github.com/dongzhaohe321418-lab/crossaudit · MIT licence · paper: CrossAudit (arXiv preprint)